

Anh Dung CONG

LinkedIn | GitHub | Portfolio

SUMMARY

AI/ML Engineer with 2+ years of experience bridging research and production — from fine-tuning multilingual models (Vietnamese–Japanese–English) and experimenting with LoRA/QLoRA optimization, to deploying scalable systems such as an LLM Gateway orchestrating multiple AI providers and an AI-powered language learning platform. Experienced with the full ML lifecycle: data pipelines, model training, API design, and CI/CD deployment with Docker and Kubernetes. Driven by a passion for building AI solutions that are both technically rigorous and practically impactful.

EDUCATION

VNU University of Science (HUS)

Computer and Information Science

Hanoi, Vietnam

Mar 2019 – Jun 2025

- Relevant coursework: Machine Learning, Natural Language Processing, Algorithms & Data Structures, Computer Vision

EXPERIENCE

AI Engineer

eUp Group

Jan 2025 – Present

Hanoi, Vietnam

Mazii Dictionary

Researched and implemented a recommendation system to personalize content on Mazii's community page, improving user engagement and retention.

Marketing Push Notification Pipeline

Designed and deployed an end-to-end AI-driven push notification pipeline: generating notification content from user descriptions, processing large-scale user data, performing real-time user segmentation based on tracking data, and delivering notifications to devices via device ID.

Market Intelligence and Data Analytics Pipeline

Built a data pipeline for marketing market analysis: crawling multi-source data, annotating and preprocessing datasets, applying AI-based insights extraction, and delivering actionable analytical reports for business decision-making.

Japanese ASR and MLOps System

Researched, fine-tuned, and deployed Japanese ASR models; designed a full MLOps pipeline (data collection, training, evaluation, deployment, monitoring) to support speech scoring features in Japanese learning applications.

AI-generated Course Builder (Hey Apps)

Developed a generative AI system to create personalized language-learning courses by topic, enabling users to dynamically construct customized learning paths within Hey-series applications.

AI Engineer

AMELA Technology

Nov 2024 – Jan 2025

Hanoi, Vietnam

Researched OCR solutions for Japanese handwriting recognition, evaluating and benchmarking multiple deep learning approaches.

Designed and prototyped a Retrieval-Augmented Generation (RAG) pipeline for an enterprise chatbot system.

Investigated AI-based automated cosmetic eyeliner application using computer vision and generative models.

AI Engineer Intern

FPT Smart Cloud

Mar 2023 – Nov 2024

Hanoi, Vietnam

FPT AI Enhance for Home Credit, FE Credit, MB Bank

Participated in the development and deployment of the FPT AI Enhance project.

Handled big data processing, and trained and evaluated models for the NLP task of automatic call center scoring.

FPT AI Enhance for FPT Long Chau, FPT Shop

Built and deployed the FPT AI Enhance pipeline.

Managed big data processing, model training and evaluation, and automation of processes and data integration using n8n.

Guardrails for LLMs

Developed and tested the Guardrails project to improve the safety and control of large language models (LLMs).

Virtual Assistant Chatbot for Various Organizations

Contributed to developing and deploying virtual assistant chatbot projects for organizations such as the State Securities Commission of Vietnam, Home Credit, and FE Credit.

PROJECTS

AI Gateway | [GitHub](#)
Personal Project

2025
Java, Spring Boot WebFlux, Python, FastAPI, Docker, Redis

Built a production-grade AI Gateway combining Java (Spring Boot WebFlux) gateway and Python (FastAPI) worker for unified AI request management.

Implemented multi-provider routing (OpenAI, Gemini, Claude, local LLM) with fallback, retry logic, circuit breaker, and rate limiting.

Developed a Python worker supporting local LLM inference, embeddings, and a RAG pipeline with document indexing and search.

Integrated Redis caching, Prometheus monitoring, and deployed the full system via Docker Compose.

Research Agent | [GitHub](#)
Personal Project

2025
Python, LangChain, LangGraph, Gemini API, MCP, A2A Protocol

Built an AI-powered research assistant for automated literature research, paper discovery, and report generation.

Designed a multi-agent architecture (Search, Paper, Synthesis agents) communicating via A2A Protocol with MCP Gateway integration.

Implemented multi-source search (Web + ArXiv), iterative research expansion, paper summarization, and structured Markdown report generation with citations.

Exposed as an MCP server for integration with Claude Desktop, Cursor, and other MCP clients.

Whisper Japanese ASR Finetune | [GitHub](#)
Personal Project

2025
Python, PyTorch, HuggingFace, LoRA, Kaggle, GitHub Actions

Fine-tuned OpenAI Whisper Tiny for Japanese speech recognition using LoRA on ReazonSpeech dataset.

Built full MLOps pipeline with CI/CT/CD using GitHub Actions, Kaggle for training, and HuggingFace Hub for model hosting.

Implemented quality gate for model promotion and exported to faster-whisper (INT8) for efficient inference.

TECHNICAL SKILLS

Programming: Python, Java, Scala, Big Data (Spark), Cloud (GCP)

Libraries/Frameworks: TensorFlow, PyTorch, OpenCV, NLTK, LangChain, LangGraph, FastAPI, Spring Boot WebFlux

Tools: Git, Linux, Docker, Redis, MongoDB, n8n, HuggingFace, GitHub Actions, Prometheus, Kaggle

English: B1 – able to communicate and read technical documents

Soft Skills: Communication, Presentation, Teamwork, Research & Paper Implementation

CERTIFICATIONS

Mathematics for Machine Learning and Data Science Specialization (2023)

Machine Learning Specialization (2023)

HONORS & AWARDS

Second Prize in Scientific Research Project (Faculty level) – May 2022

Third Prize in Scientific Research Project (Faculty level) – May 2024

REFERENCES
